


DOCKER DEEP-DIVE

Understanding how Docker leverage Linux features and its layer-based image mechanism
to create the magic container world.



Lu Hoang Anh - DevOps Engineer
Application Architecture - Cloud Architecture Team



Containers



Images



Networking



Storage

Agenda

A. Docker - VM Killer

1. What is Docker?
2. Why we need Docker?
3. Why we call it “VM Killer”
4. Docker Architecture
5. Docker Terminology

B. Docker deep-dive

C. Lesson learned & Wrap-up

DOCKER MARKET OVERVIEW

65%

use Docker to deliver development agility.

48%


use Docker to control app environments.

41%

use Docker to achieve app portability.

90%


use Docker for apps in development.



90%

plan dev environments around Docker.

Docker Workloads



58%

use Docker for apps in production.



80%

plan DevOps around Docker.

A. Docker -
VM Killer



A. Docker - VM Killer

1. What is Docker?



Docker is an open platform for deploying applications in **lightweight, portable** software containers.

Each container includes:

-  Application code
-  Libraries & dependencies
-  System tools

Docker ensures applications run **the same way** on any machine.

A. Docker - VM Killer

1. What is Docker?



Docker is an open platform for deploying applications in **lightweight, portable** software containers.

Each container includes:

-  Application code
-  Libraries & dependencies
-  System tools

Docker ensures applications run **the same way** on any machine.

2. Why do we need Docker?

Faster delivery

Standardized development, testing, and production environments accelerate software delivery.

Convenient packaging

Package applications and dependencies into a single unit for simplified deployment.

Cross-environment consistency

Solves "it works on my machine" problems with consistent runtime environments.

Easy and clear monitoring

Provides a unified way to read log files from all running containers. No need to remember all the specific paths

Scalability

Containerized applications can be easily scaled horizontally to handle growing loads.

A. Docker - VM Killer

3. Why we call it “VM Killer”

What is the difference between VMs and Docker Container?

A. Docker - VM Killer

3. Why we call it “VM Killer”



Virtual Machines



Complete OS

Each VM includes a complete operating system




Resource Heavy

High resource consumption due to multiple OS instances



Slow Startup

Time-consuming boot process



VS



Docker Containers



Shared OS

Containers share the host machine's operating system kernel




Lightweight

Lower resource consumption, more efficient use of system resources




Fast Startup

Containers can be started in seconds




A. Docker - VM Killer

3. Why we call it “VM Killer”




A. Docker - VM Killer

3. Why we call it “VM Killer”




7 Gb of Memory is blocked and cannot be allotted to a new VM



Only 9 GB memory utilized;
7 GB can be allotted to a new Container

A. Docker - VM Killer


4. Docker Architecture



Docker architecture - Source: <https://docs.docker.com/get-started/docker-overview/>

A. Docker - VM Killer

5. Docker Terminology



A. Docker - VM Killer

5. Docker Terminology



Container

A running instance of an application, including the application and all its dependencies.



Image

A read-only template used to create containers, containing application code, runtime, libraries, and configuration.



Port

Used to expose container-internal services to the external network, enabling communication.



Volume

Used to persist container data so it can be preserved beyond the container's lifecycle.



Network

Allows communication between containers and with the external world.



Registry

A service for storing and distributing Docker images, such as Docker Hub.



B. Docker Deep-dive

B. Storytelling: Docker Concepts in-action


The Project Story



A simple **todo-app** which will have 2 components: logic and database

What are the possible problems when we start developing this app?

- </> Developers are using different environments (Windows, Mac, Linux)
- 📦 Manually setting up the environment is time-consuming and error-prone
- ✗ Famous problem: "It works on my machine"



B. Storytelling: Docker Concepts in-action

1. Understand Docker Images & Containers

Docker Containers

The smallest executable software unit, packaging all needed parts to run an application.

Key characteristics:


- Runs as a process on the host machine
- Uses Linux kernel features for isolation
- Ephemeral (no persistent state)

Container Isolation

cgroups
Limit resource usage

chroot
Changes root directory

namespaces
Isolates users/processes



Source: <https://blog.devops.dev/linux-containers-deep-dive-c0668a4f347d>

B. Storytelling: Docker Concepts in-action

1. Understand Docker Images & Containers


Docker Images

A package containing all files needed to create containers.

Copy-on-write modal: When create an image, every step is cached and can be reused in future builds.

Two important principles:

- Images are Immutable - changes require new image
- Images are composed of layers



The screenshot shows the Docker Hub page for the official Python image. The URL is https://hub.docker.com/_/python. The page includes the Docker logo, a search bar with "python", and navigation links for "Explore", "Official Images", and "python". The main content area displays the "python" image card, which features the Python logo, the text "Docker Official Image", and "Python is an interpreted, interactive, object-oriented, open-source programming language." Below the card are sections for "Overview" and "Tags". To the right, there's a "Quick reference" section with links to maintainers and help resources, and a "Supported tags and respective Dockerfile links" section listing various tags like "3.14.0-trixie", "3.14-trixie", etc. On the far right, there are "Tag summary", "Content type", "Digest", "Size", and "Last updated" details, along with a "Run in Docker Desktop" button.

Source: https://hub.docker.com/_/python

B. Storytelling: Docker Concepts in-action


1. Understand Docker Images & Containers

Image Layers

- A layer is essentially a **snapshot / diff** of the filesystem captured at one point
- Layers are content-addressable
- Layers are stacked on top of each other to form the final image.
- Layers are cached and reused to optimize build times and reduce storage usage.

Example: a theoretical image

1. The first layer adds basic commands and a package manager, such as apt.
2. The second layer installs a Python runtime and pip for dependency management.
3. The third layer copies in an application's specific requirements.txt file.
4. The fourth layer installs that application's specific dependencies.
5. The fifth layer copies in the actual source code of the application.



Source: <https://docs.docker.com/get-started/docker-concepts/building-images/understanding-image-layers/>

B. Storytelling: Docker Concepts in-action


1. Understand Docker Images & Containers

Image Layers

- A layer is essentially a **snapshot / diff** of the filesystem captured at one point
- Layers are content-addressable
- Layers are stacked on top of each other to form the final image.
- Layers are cached and reused to optimize build times and reduce storage usage.

Example: create another Python application

1. Leverage the same Python base
2. Make builds faster
3. Reduce the amount of storage and bandwidth



Source: <https://docs.docker.com/get-started/docker-concepts/building-images/understanding-image-layers/>

B. Storytelling: Docker Concepts in-action

1. Understand Docker Images & Containers


Image Layers

How layers are stacked to form an image

1. Image = layer digests + image manifest (metadata)
2. When run a container, an extra writable layer is added on the top
→ changes happened inside container are isolated in this writable layer
3. At runtime, Docker mounts these read-only layers using a **union filesystem** → single merged view

What is the benefit?

Union filesystem terminology



Source: <https://docs.docker.com/get-started/docker-concepts/building-images/understanding-image-layers/>

B. Storytelling: Docker Concepts in-action

1. Understand Docker Images & Containers


Image Layers

How layers are stacked to form an image

1. Image = layer digests + image manifest (metadata)
2. When run a container, an extra writable layer is added on the top
→ changes happened inside container are isolated in this writable layer
3. At runtime, Docker mounts these read-only layers using a **union filesystem** → single merged view

What is the benefit?

Union filesystem terminology



Source: <https://docs.docker.com/get-started/docker-concepts/building-images/understanding-image-layers/>



"How do we create images for our app? And more importantly, how do we define what goes into the image?"

B. Storytelling: Docker Concepts in-action

2. Design Dockerfile

What is Dockerfile?

- A Dockerfile is a text-based document that's used to create a container image
- Provide instructions to the image builder, e.g. run cmd, copy file, set env variable,...

Common Instructions

FROM Base image

WORKDIR Set working directory

COPY Copy files from host to image

ADD Add local or remote files and directories

RUN Execute commands in image

CMD Default command to run

ENTRYPOINT Specify default executable

Dockerfile Example

```
FROM python:3.13           # Base image
WORKDIR /usr/local/app     # Set working directory

# Install dependencies
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt      # Run command inside image

# Copy in the source code
COPY src ./src
EXPOSE 8000               # Document the port the app listens on

# Copy in the source code
RUN useradd app_user        # Set default user for all subsequent instructions
USER app

CMD unicorn app.main:app \
    --host 0.0.0.0 \
    --port 8080              # Default command to run when container starts
```

B. Storytelling: Docker Concepts in-action


3. Docker Storage



What is Docker Volume?

Docker Volume provide the ability to:

- Connect specific filesystem paths of the container back to the host machine
- Allow data to persist beyond the lifecycle of a container.



B. Storytelling: Docker Concepts in-action

2. Docker Storage



Volume Mounts

Managed by Docker, persistent data storage.

- ✓ Managed by Docker

- ✓ Stored in Docker's storage area

```
>_ docker volume create mydata
```



Bind Mounts

Direct connection between host and container.


- ✓ Direct mapping of host paths

- ✓ Good for development and sharing

```
>_ docker run -v /host/path:/container/path
```

B. Storytelling: Docker Concepts in-action

2. Docker Storage



Source: <https://therahulsarkar.medium.com/understanding-docker-volumes-a-comprehensive-guide-46339aa9ac53>

B. Storytelling: Docker Concepts in-action


4. Docker Networking



What is Docker Network?

A Docker Network is a **virtual network** that allows isolated containers to communicate with each other and with the outside world.

Docker Network Types





Source: <https://www.geeksforgeeks.org/devops/basics-of-docker-networking/>

B. Storytelling: Docker Concepts in-action

➡ Bridge (Default)

Single host container default network.

- ✓ Containers can communicate
- ✓ Isolated from other networks



Bridge Type

Source: <https://www.geeksforgeeks.org/devops/basics-of-docker-networking/>



Docker Networking Fundamentals

4. Docker Networking

Host

Shares host machine network stack.

- ✓ High performance
- ⚠ No isolation



Host Type

Source: <https://www.geeksforgeeks.org/devops/basics-of-docker-networking/>

B. Storytelling: Docker Concepts in-action



4. Docker Networking



None

No network connection.

✓ Complete isolation



None Type

Source: <https://www.geeksforgeeks.org/devops/basics-of-docker-networking/>

B. Storytelling: Docker Concepts in-action


5. Docker Compose

What is Docker Compose?

- Tool for defining and running multi-container Docker applications
- Configuration via YAML file
- CLI commands to start/stop all services

Why Use Docker Compose?

- Easy management:** Simplifies deployment and management of complex applications
- Enhanced collaboration:** Team members can work in the same environment through shared Compose files
- Time savings:** Reuse existing container configurations
- Cross-environment compatibility:** Support for different settings (configuration profiles, override mechanisms)



Source: https://hitolog.blog/2022/02/19/docker_tutorial_docker-compose/

B. Storytelling: Docker Concepts in-action

6. Docker Security



B. Storytelling: Docker Concepts in-action

6. Docker Security



Common threats






Image vulnerabilities



Container escape



Data breaches



Insecure configuration

B. Storytelling: Docker Concepts in-action

6. Docker Security

🛡️ Security practices

👤 Principle of Least Privilege

Containers and users should have only the minimum privileges needed.

✅ Official Images

Use official images that are typically more secure and regularly updated.

👤 Non-root Users

Run containers as non-root users to reduce potential damage.

🚫 Limit Container Capabilities

Restrict system calls that containers can execute.

🔍 Image Vulnerability Scanning

Use tools to detect known vulnerabilities in images.

⟳ Regular Updates

Keep Docker and images up to date with security patches.


⚙️ Production Readiness

📊 Performance Optimization

- ✓ Multi-stage builds to reduce image size
- ✓ Optimized layer caching for build times
- ✓ Resource limits for CPU and memory

🛠️ Troubleshooting & Debugging

- ✓ Docker logs for container output
- ✓ Docker exec to run commands in containers
- ✓ Health checks for container status



C. Lesson Learned & Wrap-up

C. Lesson Learned & Wrap-up

Recap key points


- ✓ Docker basics: container, image, Dockerfile, networking, volume
- ✓ Docker Compose for multi-container apps
- ✓ Docker security best practices
- ✓ Production readiness: performance optimization, troubleshooting

Lesson learned

- ✓ Docker ensures consistent environments → no “works on my machine”.
- ✓ Images, containers, networks, and volumes are the building blocks.
- ✓ Compose helps manage multi-container apps easily.
- ✓ Security and optimization are essential in production.

Advanced topics


- ✓ Orchestration with Kubernetes
- ✓ CI/CD integration
- ✓ Docker Swarm
- ✓ Docker in cloud environments (AWS, GCP, Azure)



Docker, Docker



Yes papa



Eating RAM ?



No papa



Telling lies ?



No papa




Open your mouth!

Process Name	Memory
com.docker.hyperkit	9.06 GB

Reference

- <https://docs.docker.com/get-started/workshop/>
- <https://github.com/alexryabtsev/docker-workshop>
- <https://docs.docker.com/>
- <https://dev.to/nobleman97/docker-networking-101-a-blueprint-for-seamless-container-connectivity-3i5b>
- <https://blog.devops.dev/linux-containers-deep-dive-c0668a4f347d>



Thank you !

